

# NoSQL

Prepared By,

Tirth Shah (150410116107)  
Vrajesh Shah (150410116108)  
Shruti Dhuri (150410116109)  
Shubha Kanetkar (150410116110)

## What is NoSQL?

- ▶ Stands for **Not Only SQL**
- ▶ Class of non-relational data storage systems
- ▶ Usually do not require a fixed table schema nor do they use the concept of joins
- ▶ All NoSQL offerings relax one or more of the ACID properties (will talk about the CAP theorem)

## Why NoSQL?

- ▶ For data storage, an RDBMS cannot be the be-all/end-all
- ▶ Just as there are different programming languages, need to have other data storage tools in the toolbox
- ▶ A NoSQL solution is more acceptable to a client now than even a year ago
  - Think about proposing a Ruby/Rails or Groovy/Grails solution now versus a couple of years ago

## How did we get here?

- ▶ Explosion of social media sites (Facebook, Twitter) with large data needs
- ▶ Rise of cloud-based solutions such as Amazon S3 (simple storage solution)
- ▶ Just as moving to dynamically-typed languages (Ruby/Groovy), a shift to dynamically-typed data with frequent schema changes
- ▶ Open-source community

## CAP Theorem

- ▶ Three properties of a system: consistency, availability and partitions
- ▶ You can have at most two of these three properties for any shared-data system
- ▶ To scale out, you have to partition. That leaves either consistency or availability to choose from
  - In almost all cases, you would choose availability over consistency

## Consistency Model

- ▶ A consistency model determines rules for visibility and apparent order of updates.
- ▶ For example:
  - Row X is replicated on nodes M and N
  - Client A writes row X to node N
  - Some period of time t elapses.
  - Client B reads row X from node M
  - Does client B see the write from client A?
  - Consistency is a continuum with tradeoffs
  - For NoSQL, the answer would be: maybe
  - CAP Theorem states: Strict Consistency can't be achieved at the same time as availability and partition-tolerance.

## Eventual Consistency

- ▶ When no updates occur for a long period of time, eventually all updates will propagate through the system and all the nodes will be consistent
- ▶ For a given accepted update and a given node, eventually either the update reaches the node or the node is removed from service
- ▶ Known as BASE (**B**asically **A**vailable, **S**oft state, **E**ventual consistency), as opposed to ACID

## What kinds of NoSQL

- ▶ NoSQL solutions fall into two major areas:
  - Key/Value or 'the big hash table'.
    - Amazon S3 (Dynamo)
    - Voldemort
    - Scalaris
  - Schema-less which comes in multiple flavors, column-based, document-based or graph-based.
    - Cassandra (column-based)
    - CouchDB (document-based)
    - Neo4J (graph-based)
    - HBase (column-based)

## Key/Value

### *Pros:*

- very fast
- very scalable
- simple model
- able to distribute horizontally

### *Cons:*

- many data structures (objects) can't be easily modeled as key value pairs

## Schema-Less

### *Pros:*

- Schema-less data model is richer than key/value pairs
- eventual consistency
- many are distributed
- still provide excellent performance and scalability

### *Cons:*

- typically no ACID transactions or joins

## Common Advantages

- ▶ Cheap, easy to implement (open source)
- ▶ Data are replicated to multiple nodes (therefore identical and fault-tolerant) and can be partitioned
  - Down nodes easily replaced
  - No single point of failure
- ▶ Easy to distribute
- ▶ Don't require a schema
- ▶ Can scale up and down
- ▶ Relax the data consistency requirement (CAP)

THANK YOU